

Ali Nawab

214-597-1122 | mxn210071@utdallas.edu | github.com/yaboiali | bit.ly/Alis-Portfolio | linkedin.com/in/alinawab

EDUCATION

University of Texas at Dallas

Graduation: December 2027

Bachelor of Science, Computer Science

GPA: 3.7/4.0

Relevant Coursework: Data Structures, Algorithms, Software Engineering, Deep Learning, Machine Learning, Natural Language Processing, Artificial Intelligence, Statistics, Computer Architecture, Distributed Systems

Technical Skills: Python (Intermediate), Java (Intermediate), C++ (Beginner), C# (Beginner), JavaScript (Intermediate)

Others: React.js, Node.js, Firebase, Express, GraphQL, Flask, Spring boot, ASP.NET, Flutter, Jenkins, Spark

EXPERIENCE

Bloomberg | Tech Insights Program

June 2025 – August 2025

- **Engaged in a selective**, immersive Bloomberg Tech Insights Program, gaining firsthand exposure to the Bloomberg Engineering culture and strategic approach to technical interviews.
- **Completed** weekly coaching sessions led by **Bloomberg engineers**, sharpening coding techniques, enhancing problem-solving proficiency, and refining communication skills in technical contexts.
- **Acquired structured preparation in technical interview methodology**, elevating readiness for engineering challenges through consistent practice and constructive feedback.

Palantir | Software engineer intern

February 2025 – March 2025

- Selected for Palantir's Launch program and tasked with showcasing AIP's capabilities by building a **Sentimental Gmail Organizer**—an AI-driven tool that classifies and prioritizes high-volume email streams based on real-time sentiment analysis.
- **Designed** and implemented a fully-asynchronous pipeline in AIP (**Foundry**) that ingests Gmail API data, applies large-language-model sentiment scoring, and auto-labels & routes critical messages—cutting inbox triage time.
- **Integrated** AIP ontology & security primitives, deploying to an internal pilot across five teams. Achieved **95% precision** in positive/negative classification and reduced missed high-priority emails.
- **Collaborated** with product and infra engineers to templatize the workflow, enabling rapid replication for other communication channels (Slack, Teams) and demonstrating **AIP's** extensibility to leadership.

Meta | Software engineer Intern - MLH Fellow

June 2024 – September 2024

- Selected as 1 of 50 students from a 5,000+ applicant pool to collaborate with Meta engineers, contributing to scalable AI inferencing pipelines.
- Containerized ML workloads using **Docker** and **Kubernetes**, fostering robust AI integration in production environments.
- Optimized **GPU** resource scheduling and fine-tuned microservices, reducing system downtime by **5%** and supporting high-performance system operations.
- Streamlined **CI/CD** processes by automating over **600+** pull request verifications with Jenkins pipelines, enhancing automated testing and deployment workflows

PROJECTS

Lotus AI: github.com/AnSingh1/nurseai — An agentic virtual nurse that empathetically guides users through health concerns

- **Co-built** Lotus AI, an agentic virtual nurse that performs real-time triage using multimodal inputs (voice + video) and structured reasoning.
- **Engineered** the intelligent search and context-tracking system enabling Lotus to remember symptoms, interpret user inputs, perform location-based hospital search, and trigger follow-up actions.
- **Developed** backend workflow for real-time clinical summarization, appointment booking automation, and seamless integration of speech, vision, and reasoning models.

Meddy: Back-end Engineer – <https://meddy-kqfo.vercel.app/>

- Led backend architecture and development for an AI-driven healthcare assistant delivering voice-based medical insights, treatment reminders, and real-time Q&A.
- Designed microservice-based architecture using **AWS EC2** for compute, **Lambda** for serverless function execution, and **DynamoDB** for low-latency **NoSQL storage**.
- Integrated **Gemini API** for LLM-powered medical responses and **Firebase** for auth and secure data sync.

AI Research: Llama.cpp Integration

January 2024 – March 2024

- Researched MetaAI's **LLaMA2** to assess its **Natural Language Processing** capabilities, evaluating prompt engineering quality and response performance through analytical metrics
- Trained and integrated an **LLM** on the CS Genome Organizations website, automating question answering of CS Genome processor descriptions and general **HPC** computing topics (**5-15 sec.** response time with **85%+ accuracy**)

LEADERSHIP

Valley Ranch Community Center | Computer Science Tutor

June 2023 - Present

- Designed and taught a weekly CS curriculum (Python, web dev, intro ML) for **25 +** high-school students, lifting average project-assessment scores from 70 % to 92 % in one term.
- Hosted quarterly hackathons and outreach workshops, leading teams to 3 regional competition wins and securing \$4 k in grants for new lab hardware.